

```
In [52]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [53]: data_train = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data', header = None)
data_test = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.test', skiprows = 1, head
```

```
In [54]: data_train.tail(3)
```

Out[54]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	0	0	40	United-States	<=50K
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	0	20	United-States	<=50K
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	0	40	United-States	>50K

```
In [55]: data_test.tail(3)
```

Out[55]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
16278	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	0	50	United-States	<=50K.
16279	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander	Male	5455	0	40	United-States	<=50K.
16280	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	60	United-States	>50K.

## Adding colum to dataset

```
In [56]: column = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status',
                  'occupation', 'relationship', 'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week',
                  'native_country', 'wage_class']
data_train.columns = column
data_test.columns = column
```

```
In [57]: data_train.head(3)
```

Out[57]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_w
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	

```
In [58]: df = pd.concat([data_train, data_test])
df['workclass'].value_counts()
```

```
Out[58]: Private      33906
Self-emp-not-inc    3862
Local-gov          3136
?                  2799
State-gov           1981
Self-emp-inc        1695
Federal-gov         1432
Without-pay         21
Never-worked        10
Name: workclass, dtype: int64
```

```
In [59]: for column in df.columns:
         print(f" value count for {column} : \n {df[column].value_counts()}")
```

```
208174      1
Name: fnlwt, Length: 28523, dtype: int64
value count for education :
HS-grad      15784
Some-college 10878
Bachelors    8025
Masters       2657
Assoc-voc     2061
11th          1812
Assoc-acdm    1601
10th          1389
7th-8th        955
Prof-school    834
9th            756
12th           657
Doctorate      594
5th-6th        509
1st-4th        247
Preschool       83
Name: education, dtype: int64
```

```
df.info()
```

```
In [60]: df.describe()
```

```
Out[60]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
<b>count</b>	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.000000
<b>mean</b>	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.422382
<b>std</b>	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.391444
<b>min</b>	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
<b>25%</b>	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.000000
<b>50%</b>	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.000000
<b>75%</b>	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.000000
<b>max</b>	90.000000	1.490400e+06	16.000000	99999.000000	4356.000000	99.000000

## replacing ? from workclass column

```
In [61]: df.replace('?',np.nan,inplace=True)
```


```
In [62]: df.wage_class.unique()
```

```
Out[62]: array([' <=50K', ' >50K', ' <=50K.', ' >50K.'], dtype=object)
```

```
In [66]: df.replace({' <=50K':0, ' >50K':1, ' <=50K.':0, ' >50K.':1}).head(3)
```

Out[66]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_w
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	




```
In [67]: df['workclass'].fillna('0',inplace=True)
```

```
In [69]: plt.figure(figsize=(10,8))  
sns.countplot(df['workclass'])  
plt.xticks(rotation = 45)  
plt.show()
```

C:\Users\kants\AppData\Local\Programs\Python\Python37\lib\site-packages\seaborn\\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



```
In [70]: df['education'].value_counts()
```


```
Out[70]: HS-grad      15784
Some-college  10878
Bachelors     8025
Masters       2657
Assoc-voc     2061
11th          1812
Assoc-acdm    1601
10th          1389
7th-8th       955
Prof-school   834
9th           756
12th          657
Doctorate     594
5th-6th       509
1st-4th       247
Preschool     83
Name: education, dtype: int64
```


```
In [71]: df.columns
```

```
Out[71]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
               'marital_status', 'occupation', 'relationship', 'race', 'sex',
               'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
               'wage_class'],
              dtype='object')
```



```
In [73]: sns.catplot(x='education',y='wage_class',data=df,height=10,palette='muted',kind='bar')
plt.xticks(rotation=45)
plt.show()
```







```
In [74]: def primary(x):  
         if x in ['1st-4th', '5th-6th', '7th-8th', '9th', '10th', '11th', '12th']:  
             return 'Primary'  
         else:  
             return x
```

```
In [75]: df['education'] = df['education'].apply(primary)
```


```
In [76]: sns.catplot(x='education',y='wage_class',data=df,height=10,palette='muted',kind='bar')
plt.xticks(rotation=45)
plt.show()
```






```
In [77]: df['marital_status'].replace(' Married-AF-spouse', ' Married-civ-spouse',inplace=True)
```

```
In [78]: sns.catplot(x='marital_status',y='wage_class',data=df,palette='muted',kind='bar',height=8)
plt.xticks(rotation=45)
plt.show()
```



```
In [79]: df['occupation'].fillna('0',inplace=True)
df['occupation'].value_counts()
df['occupation'].replace(' Armed-Forces','0',inplace=True)
df['occupation'].value_counts()
sns.catplot(x='occupation',y='wage_class',data=df,palette='muted',kind='bar',height=8)
plt.xticks(rotation=45)
```

```
Out[79]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]),
 [Text(0, 0, ' Adm-clerical'),
  Text(1, 0, ' Exec-managerial'),
  Text(2, 0, ' Handlers-cleaners'),
  Text(3, 0, ' Prof-specialty'),
  Text(4, 0, ' Other-service'),
  Text(5, 0, ' Sales'),
  Text(6, 0, ' Craft-repair'),
  Text(7, 0, ' Transport-moving'),
  Text(8, 0, ' Farming-fishing'),
  Text(9, 0, ' Machine-op-inspct'),
  Text(10, 0, ' Tech-support'),
  Text(11, 0, ' ?'),
  Text(12, 0, ' Protective-serv'),
  Text(13, 0, '0'),
  Text(14, 0, ' Priv-house-serv')])
```





```
In [81]: df['relationship'].value_counts()
```

```
Out[81]: Husband          19716  
Not-in-family    12583  
Own-child        7581  
Unmarried        5125  
Wife             2331  
Other-relative   1506  
Name: relationship, dtype: int64
```

```
In [82]: df['race'].value_counts()
```

```
Out[82]: White           41762  
Black           4685  
Asian-Pac-Islander   1519  
Amer-Indian-Eskimo    470  
Other             406  
Name: race, dtype: int64
```

```
In [83]: df.columns
```

```
Out[83]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',  
               'marital_status', 'occupation', 'relationship', 'race', 'sex',  
               'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',  
               'wage_class'],  
              dtype='object')
```

```
In [84]: df['sex'].value_counts()
```

```
Out[84]: Male      32650  
Female    16192  
Name: sex, dtype: int64
```

```
In [85]: df['native_country'].unique()
```


```
Out[85]: array([' United-States', ' Cuba', ' Jamaica', ' India', ' ?', ' Mexico',  
              ' South', ' Puerto-Rico', ' Honduras', ' England', ' Canada',  
              ' Germany', ' Iran', ' Philippines', ' Italy', ' Poland',  
              ' Columbia', ' Cambodia', ' Thailand', ' Ecuador', ' Laos',  
              ' Taiwan', ' Haiti', ' Portugal', ' Dominican-Republic',  
              ' El-Salvador', ' France', ' Guatemala', ' China', ' Japan',  
              ' Yugoslavia', ' Peru', ' Outlying-US(Guam-USVI-etc)', ' Scotland',  
              ' Trinidad&Tobago', ' Greece', ' Nicaragua', ' Vietnam', ' Hong',  
              ' Ireland', ' Hungary', ' Holand-Netherlands'], dtype=object)
```

```
In [86]: def native(country):  
    if country in [' United-States', ' Canada']:  
        return 'North_America'  
    elif country in [' Puerto-Rico', ' El-Salvador', ' Cuba', ' Jamaica', ' Dominican-Republic', ' Guatemala', ' Haiti', ' Nic  
        return 'Central_America'  
    elif country in [' Mexico', ' Columbia', ' Vietnam', ' Peru', ' Ecuador', ' South', ' Outlying-US(Guam-USVI-etc)']:  
        return 'South_America'  
    elif country in [' Germany', ' England', ' Italy', ' Poland', ' Portugal', ' Greece', ' Yugoslavia', ' France', ' Ireland',  
        return 'EU'  
    elif country in [' India', ' Iran', ' China', ' Japan', ' Thailand', ' Hong', ' Cambodia', ' Laos', ' Philippines', ' Taiwan  
        return 'Asian'  
    else:  
        return country
```

```
In [88]: df['native_country'] = df['native_country'].apply(native)
```


```
In [89]: sns.catplot(x='native_country',y='wage_class',data=df,palette='muted',kind='bar',height=8)
plt.xticks(rotation=45)
```

```
Out[89]: (array([0, 1, 2, 3, 4, 5]),
 [Text(0, 0, 'North_America'),
  Text(1, 0, 'Central_America'),
  Text(2, 0, 'Asian'),
  Text(3, 0, ' '),
  Text(4, 0, 'South_America'),
  Text(5, 0, 'EU')])
```



```
In [90]: corr = df.corr()  
plt.figure(figsize=(10,12))  
sns.heatmap(corr,annot=True)
```

Out[90]: <AxesSubplot:>





```
In [91]: X = df.drop(['wage_class'],axis=1)
        y = df['wage_class']
```

```
In [92]: X.columns
```

```
Out[92]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
               'marital_status', 'occupation', 'relationship', 'race', 'sex',
               'capital_gain', 'capital_loss', 'hours_per_week', 'native_country'],
              dtype='object')
```


```
In [94]: X_dummy = pd.get_dummies(X)
```

```
In [95]: X_dummy.head()
```

Out[95]:

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week	workclass_?	workclass_Federal-gov	workclass_Local-gov	workclass_Never-worked	...	race_Other	race_White	Fe
0	39	77516	13	2174	0	40	0	0	0	0	...	0	1	
1	50	83311	13	0	0	13	0	0	0	0	...	0	1	
2	38	215646	9	0	0	40	0	0	0	0	...	0	1	
3	53	234721	7	0	0	40	0	0	0	0	...	0	0	
4	28	338409	13	0	0	40	0	0	0	0	...	0	0	

5 rows × 65 columns



```
In [96]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_d)
```

```
In [97]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X_scaled,y,test_size=0.3,random_state=101)
```

```
In [98]: parameters = [{ 'learning_rate':[0.01,0.001],
                        'max_depth': [3,5,10],
                        'n_estimators':[10,50,100,200]
                      }
                    ]
```

```
In [102]: from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier
Xbc = XGBClassifier()
Grid_cv = GridSearchCV(Xbc,parameters,scoring='accuracy',cv=5,n_jobs=3,verbose=3)
Grid_cv.fit(x_train,y_train)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

C:\Users\kants\AppData\Local\Programs\Python\Python37\lib\site-packages\xgboost\sklearn.py:1146: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use\_label\_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num\_class - 1].

warnings.warn(label\_encoder\_deprecation\_msg, UserWarning)

[14:33:37] WARNING: C:/Users/Administrator/workspace/xgboost-win64\_release\_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval\_metric if you'd like to restore the old behavior.

```
Out[102]: GridSearchCV(cv=5,
                    estimator=XGBClassifier(base_score=None, booster=None,
                                            colsample_bylevel=None,
                                            colsample_bynode=None,
                                            colsample_bytree=None, gamma=None,
                                            gpu_id=None, importance_type='gain',
                                            interaction_constraints=None,
                                            learning_rate=None, max_delta_step=None,
                                            max_depth=None, min_child_weight=None,
                                            missing=nan, monotone_constraints=None,
                                            n_estimators=100, n_jobs=None,
                                            num_parallel_tree=None, random_state=None,
                                            reg_alpha=None, reg_lambda=None,
                                            scale_pos_weight=None, subsample=None,
                                            tree_method=None, validate_parameters=None,
                                            verbosity=None),
                    n_jobs=3,
                    param_grid=[{'learning_rate': [0.01, 0.001],
                                'max_depth': [3, 5, 10],
                                'n_estimators': [10, 50, 100, 200]}],
                    scoring='accuracy', verbose=3)
```



```
In [104]: Grid_cv.best_params_
```

```
Out[104]: {'learning_rate': 0.01, 'max_depth': 10, 'n_estimators': 200}
```

```
In [105]: XBC = XGBClassifier(learning_rate=0.01,max_depth=10,n_estimators=200)
XBC.fit(x_train,y_train)
```

[14:38:05] WARNING: C:/Users/Administrator/workspace/xgboost-win64\_release\_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval\_metric if you'd like to restore the old behavior.

```
Out[105]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.01, max_delta_step=0, max_depth=10,
                        min_child_weight=1, missing=nan, monotone_constraints='()',
                        n_estimators=200, n_jobs=8, num_parallel_tree=1, random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [106]: XBC.score(x_test,y_test)
```

```
Out[106]: 0.8658295229645806
```

```
In [108]: y_pred = XBC.predict(x_test)
```

```
In [109]: from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
```

```
In [112]: print(f'Accuracy Score:{accuracy_score(y_test,y_pred)}')
print(f'Confusion Matrix:{confusion_matrix(y_test,y_pred)}')
print(f'Classification Report: {classification_report(y_test,y_pred)}')
```

Accuracy Score:0.8658295229645806  
Confusion Matrix:[[10536 564]  
[ 1402 2151]]  
Classification Report:

			precision	recall	f1-score	support
	0	0.88	0.95	0.91	11100	
	1	0.79	0.61	0.69	3553	
	accuracy			0.87	14653	
	macro avg	0.84	0.78	0.80	14653	
	weighted avg	0.86	0.87	0.86	14653	

```
In [ ]:
```